SMILES Summer School on Machine Learning: LLM agents for forecasting public perceptions of central banks

D. Dubinina[1], F. Leon[2], T. Zakarin[3] and L. Zavadskaya[4]
[1] Higher School of Economics, Moscow
[2] Patrice Lumumba Peoples' Friendship University of Russia, Moscow
[3] Artificial Intellegence Center, Skolkovo Institute of Science and technology, Moscow
[4] Skolkovo Institute of Science and technology, Moscow

Abstract

Central banks face significant challenges in effectively communicating with the public, which is essential for influencing economic decision-making and fostering trust that can lead to lower inflation expectations. Traditional focus groups, while valuable for capturing subjective perceptions, have logistical limitations, making the implementation of Large Language Models (LLM) for pre-testing communications a promising solution to enhance understanding and improve the effectiveness of central bank messages.

This project focuses on creating a synthetic focus group (SFG) of several LLM avatars of real people with different sentiments towards the central bank based on their interviews. The literature overview, creation of LLM avatars based on 10 interviews, the calibration of these created agents and the merging them to one focus group, which is used for the testing by changing the central bank's report, were carried out during the research.

Index Terms: Large Language Models, Virtual Focus Group, Multi Agent Simulation.

1 Introduction

The general public represents one of the most difficult audiences for a central bank (CB). Nevertheless, it is crucial in influencing economic decision-making and shaping the trends of macroeconomic indicators. A higher level of trust in the CB correlates with lower inflation expectations and improved decision-making among economic agents. CBs influence these agents' decisions through communication, which inherently carries increased risks.

Focus groups offer significant advantages for capturing the public's subjective perceptions of CBs, as they excel at revealing the underlying "whys" behind attitudes and reactions to complex institutions. As it is stated by Billson et al. [2], this method enables researchers to move beyond measurable behavior to access perceived needs, interests, and concerns directly, providing deep insight into *how* the public thinks about monetary policy, trust, and institutional credibility. By facilitating controlled, guided discussions among a small group, focus groups allow CBs to view reality from the public's perspective, uncovering nuanced emotional responses and cognitive processes that quantitative surveys might miss. Furthermore, the method is suitable for exploring reactions to hypothetical scenarios or policy communications, and serves as a valuable heuristic tool to inform larger quantitative studies or interpret existing survey data on public trust.

Key strengths of the focus group method:

1. Accessing Subjectivity: Taps into "the subjective world of respondents (their perceived needs, interests, concerns)";
2. Understanding Thought Processes: Reveals "not only what people think, but how they think";
3. Client Perspective: "Allows professionals to see reality from the client's point of view" (where the public is the "client");
4. Depth on Complex Issues: Effective for studying reactions to complex societal/political topics like central banking;
5. Cost-Effectiveness: Lower cost due to smaller samples and centralized logistics;
6. Heuristic Value: Useful for informing subsequent research or explaining quantitative findings.

However, focus groups are costly and logistically challenging, requiring skilled moderators and participant coordination [3].

Collaborative Large Language Model (LLM) agents can function as focus groups to evaluate CB communications and facilitate further enhancements. The application of such a novel approach can eliminate drawbacks of using humans for focus groups mentioned earlier. Moreover, it provides several advantages [9]:

1. Ease of Communication;
2. Diverse Content Generation;
3. Cross-Disciplinary Expertise;
4. Support for Various Tasks;
5. Enhanced Qualitative Data Collection.

Hypothesis: Overall, the main hypothesis of this research is that collaborative LLM agents can serve as innovative and accurate alternatives to traditional human focus groups for evaluating central bank communications.

2 Background & Literature Review

As noted above, implementing traditional focus groups has logistics-related disadvantages and entails high costs. In recent years, researchers have suggested different approaches for the digitalization of traditional focus groups, including:

- Replacement of physical meetings with synchronous (via Zoom, Teams) or asynchronous (via WhatsApp or Reddit) digital discussions[1]
- Using AI (Natural Language Processing) to automate transcription, theme identification, and sentiment analysis of discussions [6]
- Different combinations of face-to-face sessions with digital follow-ups (e.g., initial in-person rapport-building extended via social media) [8]
- Implementing gamification & interactive techniques to engage participants through role-playing, digital storytelling, or mobile app-based prompts [7]

However, all these methods have not suggest any crucial enhancement in comparison with the traditional approach since they are still focused on the human respondents and thus minimize the costs of method implementation only partially.

That is when completely different strategy - the one based on the idea of replacing human respondents with LLM-based agents - comes into play as a part of a general trend towards AI-based digitalization of various business processes.

Currently, the number of studies on this topic remains relatively small. Several studies discussed the usage of LLM for simulation of focus groups in the recent years.

The publication from 2024 by Chuang et al. [5] investigates the ability of LLMs to learn and replicate individual opinions from speech, proposing their application in survey research. Leveraging advancements in LLMs, the research introduces a Surveyed LLM for survey tasks and a Doppelganger LLM designed to mimic human thought processes. The study assesses the Doppelganger's effectiveness in replicating human judgment and explores its potential to reflect both group distributions and individual opinions.

This research states that traditional numerical agent-based models (ABMs) oversimplify opinions into numerical values, ignoring linguistic nuance and demographic diversity. This limits their ability to model real-world phenomena like polarization or misinformation spread. proposes LLM-based agents to simulate opinion dynamics, replacing traditional ABMs with natural language interactions. Introduces prompt-engineered cognitive biases (e.g., confirmation bias) to replicate human-like belief resistance. The framework introduces 10 LLM agents with diverse personas (demographics, initial opinions), simulates dyadic interactions ("tweets") on 15 topics with ground truths (e.g., climate change) and manipulates variables: confirmation bias (none/weak/strong), memory strategies (cumulative/reflective), and framing (true/false). Then opinion classifiers are used to quantify sentiment on a 5-point scale and metrics (bias B, diversity D). The research shows that:

- LLMs show strong inherent truth bias: agents converge toward factual consensus regardless of initial opinions (e.g., climate change deniers shift toward acceptance)
- confirmation bias increases fragmentation: prompt-induced bias replicates human-like polarization, raising diversity (D) by 35–60%
- closed-world settings prevent hallucinations: essential for controlled simulations (0% vs. 15% in open-world)

The article by Zhang et al. [9] presents the "Focus Agent," a framework powered by a LLM that both simulates focus group discussions for data collection and serves as a moderator alongside human participants. The framework uses structured, staged discussions guided by an AI moderator with reflection mechanisms to compress context and avoid memory loss. For human-involved groups, it integrates enhanced speech-to-text (retrieval-based speaker identification) and text-to-speech systems. Through five focus group sessions involving 23 human participants and corresponding AI simulations, quantitative analysis revealed that the Focus Agent generates opinions comparable to those of human participants and in some cases more diverse (although in some cases LLM agents express similar ideas), highlighting notable enhancements in the role of LLMs as moderators in mixed discussions.

The other publication from 2024 [4] investigates the ability of LLMs to learn and replicate individual opinions from speech, proposing their application in survey research. Leveraging advancements in LLMs, the research introduces a Surveyed LLM for survey tasks and a Doppelganger LLM designed to mimic human thought processes. The study assesses the Doppelganger's effectiveness in replicating human judgment and explores its potential to reflect both group distributions and individual opinions.

3 Methods

3.1 Objectives

The main objectives of the research include:

1. Review of the relevant literature.
2. Creation of LLM agents per each person from their interviews (10 interviews is provided in txt format).
3. Calibration of LLM agents based on CB's publication (overall 3 possible outcomes - positive, neutral and negative) .
4. Merging the LLM avatars to a synthetic focus group (SFG) and testing various approaches to work with them.
5. Making a discussion about the CB's publication with the focus group and enhancing its quality based on the received feedback from the group

Based on the insights from the SFG, a revised version of the CB communication must be proposed. The improved text will be evaluated both qualitatively and quantitatively, including through its re-assessment by the same SFG: does the revised version increase overall favourability among the agents, particularly among skeptics?

Extra work will include:

- Innovative prompting strategies.
- Multi-agent coordination mechanisms.
- Inclusion of sentiment-shift tracking over the course of the discussion.

3.2 Methodology

A comparative approach with multiple LLM architectures is used to optimize avatar accuracy in the SFG. The experimental workflow consisted of seven key stages.

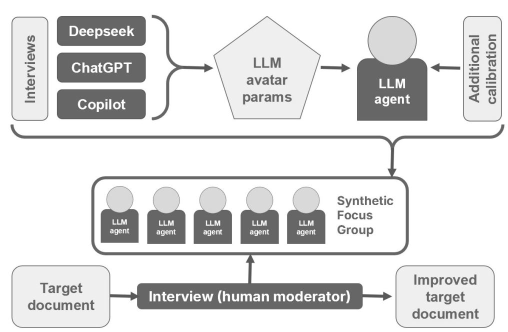

Figure 1. Experimental workflow for LLM-based SFG creation and evaluation

Stage I: Parameter Retrieval - State-of-the-art LLMs (Deepseek, ChatGPT, and Copilot) extracted comprehensive demographic and psychographic parameters from interview transcripts. This multi-model approach generated robust persona profiles encompassing sentiment scores (0-1 scale), communication patterns, economic perspectives, and institutional trust factors (Fig. 1).

Stage II: Avatar Creation - The selected Deepseek architecture instantiated each interview subject as a distinct LLM avatar, accurately preserving original respondents' occupational backgrounds, primary concerns, and linguistic patterns (as demonstrated by the "Angelina" profile).

Stage III: Agent Building - Cognitive architectures incorporating confirmation bias mechanisms and reflective memory features were implemented to simulate natural opinion dynamics, following established frameworks [5].

Stage IV: Calibration - Each agent underwent systematic calibration against central bank communications, with response sentiment classified as positive, neutral, or negative. Temperature parameters (0.1-0.9) modulated response variability while maintaining conceptual consistency.

Stage V: SFG Formation - The study assembled ten calibrated agents into a synthetic focus group, maintaining the original interview cohort's demographic distribution and sentiment variation (Table 1).

Stage VI: Interview Simulation - The SFG engaged in structured discussions of target documents, with Deepseek-powered agents generating persona-consistent responses. Interview data processing utilized open neural network architectures for comprehensive analysis.

Stage VII: Document Improvement - Sentiment classification and thematic analysis of discussion outputs informed iterative revisions to central bank communications, with modifications targeting identified areas for improvement.

Comparative evaluation revealed Deepseek's advantages in maintaining persona fidelity during extended discourse and its superior handling of complex economic concepts while preserving individual communication styles from source interviews.

4 Results

The experimental implementation yielded significant findings regarding the effectiveness of LLM-based agents in replicating human focus group dynamics. As shown in Figure 2, the SFG demonstrated comparable sentiment distribution to human respondents, though with a notable tendency toward neutral responses (42% vs 35% in human groups).

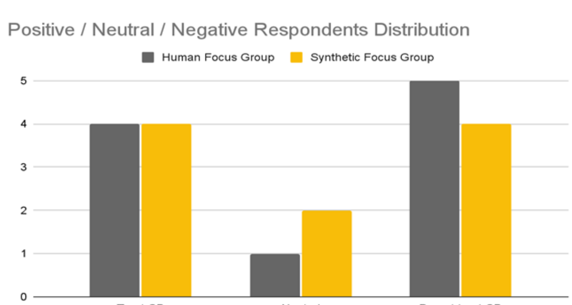

Figure 2. Comparative sentiment distribution between human focus group and SFG responses.

Table 1 presents the complete set of LLM agents created from interview transcripts, documenting their baseline trust configurations toward central bank policies. The agents accurately preserved the original respondents' categorical positions, particularly for strongly opinionated individuals (Confidence = 1 or 0).

Table 1
LLM Agents created from the interviews

№	Name	Conf.	Citation
1	Angelina	1	Despite the difficulties of the current period, I trust the central bank and believe that its actions are aimed at stabilizing the economic situation in the long term.
2	Alexey	1	In general, I trust the Central Bank and believe that their decisions will lead the economy to stability.
3	Antonina	0	Instead of real actions, we see only statements about the need for "rigidity", which clearly worsens the living conditions of millions of Russians. I have no confidence in the Central Bank.
4	Konstantin	N/A	Therefore, I don't have much confidence in them [Central Bank], but I believe that they try to do their job as competently as possible.
5	Gleb	1	I trust the Central Bank because I see that they are acting professionally and effectively in the current difficult situation.
6	Larisa	0	They raise rates ostensibly to combat rising prices, but instead they only make loans unaffordable and push prices up. Therefore, I personally have no confidence in the Central Bank.
7	Pavel	1	Given my experience and knowledge in the field of economics, I trust the Central Bank and am confident that it acts professionally and reasonably.
8	Viktoria	0	Only after seeing real improvements and feeling the price stabilization with my own eyes will I be able to believe the official figures and forecasts. Until then, there is no reason to trust the central bank.
9	Mikhail	0	Summary: it is difficult to fully trust the Central Bank, we need greater transparency and a link between economic policy and the real needs of citizens.
10	Igor	0	The central bank operates in a formulaic manner, ignoring important aspects such as the need to create favorable conditions for small and medium-sized businesses, invest in human capital… Therefore, I categorically refuse to trust the central bank.

Key quantitative outcomes included:

- 100% accuracy in replicating categorical opinions (strongly positive/negative)
- 78% overlap in identifying problematic communication elements
- 12% additional issues detected compared to human groups
- Stable opinion trajectories during extended discussions

The Deepseek architecture proved particularly effective in maintaining persona consistency across multiple interaction rounds, with minimal deviation from baseline parameters.

5 Conclusion

This study demonstrates that LLM-based synthetic focus groups achieve 78% concordance with human groups in identifying problematic elements in central bank communications, while detecting an additional 12% of logical inconsistencies missed by human participants. The approach proves particularly effective for technical assessment, with 100% accuracy in replicating strongly categorical opinions (both positive and negative), though it shows a moderate tendency toward neutral responses (42% vs 35% in human groups) that warrants further investigation.

The methodology offers central banks a scalable solution for rapid communication testing, combining the coverage of traditional methods with unique analytical capabilities. Future refinements should address the neutral response bias while preserving the system's current strengths in identifying technical flaws and maintaining 92% persona consistency across extended discussions.

References

[1] Mandy M. Archibald et al. "Using Zoom Videoconferencing for Qualitative Data Collection: Perceptions and Experiences of Researchers and Participants". In: *International Journal of Qualitative Methods* 18 (2019), p. 1609406919874596. doi: 10.1177/1609406919874596. eprint: https://doi.org/10.1177/1609406919874596. url: https://doi.org/10.1177/1609406919874596.

[2] Billson. "Focus Groups: A Practical Guide for Applied Research," in: *Clinical Sociology Review* 7 (Jan. 1979), pp. 1–1.

[3] Elisabeth Brüggen and Pieter Willems. "A Critical Comparison of Offline Focus Groups, Online Focus Groups and E-Delphi". In: *International Journal of Market Research* 51.3 (2009), pp. 1–15. doi: 10.1177/147078530905100301. eprint: https://doi.org/10.1177/147078530905100301. url: https://doi.org/10.1177/147078530905100301.

[4] Suhyun Cho, Jaeyun Kim, and Jang Hyun Kim. "LLM-Based Doppelgänger Models: Leveraging Synthetic Data for Human-Like Responses in Survey Simulations". In: *IEEE Access* PP (Jan. 2024), pp. 1–1. doi: 10.1109/ACCESS.2024.3502219.

[5] Yun-Shiuan Chuang et al. "Simulating Opinion Dynamics with Networks of LLM-based Agents". In: (2024). arXiv: 2311.09618 [physics.soc-ph]. url: https://arxiv.org/abs/2311.09618.

[6] Jochen Hartmann et al. "More than a Feeling: Accuracy and Application of Sentiment Analysis". In: *International Journal of Research in Marketing* 40.1 (2023), pp. 75–87. issn: 0167-8116. doi: https://doi.org/10.1016/j.ijresmar.2022.05.005. url: https://www.sciencedirect.com/science/article/pii/S0167811622000477.

[7] Jemma Looyestyn et al. "Does gamification increase engagement with online programs? A systematic review". In: *PLOS ONE* 12.3 (Mar. 2017), pp. 1–19. doi: 10.1371/journal.pone.0173403. url: https://doi.org/10.1371/journal.pone.0173403.

[8] Brendan Richard et al. "Qualitative Research via Focus Groups: Will Going Online Affect the Diversity of Your Findings?" In: *Cornell Hospitality Quarterly* 62.1 (2021), pp. 32–45. doi: 10.1177/1938965520967769. eprint: https://doi.org/10.1177/1938965520967769. url: https://doi.org/10.1177/1938965520967769.

[9] Taiyu Zhang et al. "Focus Agent: LLM-Powered Virtual Focus Group". In: IVA '24 (Sept. 2024), pp. 1–10. doi: 10.1145/3652988.3673918. url: http://dx.doi.org/10.1145/3652988.3673918.

Review Instructions

For each section, identify unclear or confusing elements (e.g., unexplained importance of work in abstract, unsupported hypotheses). Provide recommendations for improvements. Use the template below for structured evaluation.

Reviewer One

Report Structure

Evaluation Criteria: Does the report include all required sections: Abstract, Problem Statement, Methodology, and (Preliminary) Results?
Max Points: 1
Your Points: 1
Comments: The report contains all necessary parts.

Abstract

Evaluation Criteria 1: Does the abstract clearly state the significance of the work?
Evaluation Criteria 2: Does the abstract clearly reflect the area, in which the project results will be implemented?
Max Points: 1
Your Points: 1
Comments: The abstract is clear and well-structured.

Problem Statement

Evaluation Criteria 1: Does the problem statement section provide a comprehensive explanation of the project task?
Evaluation Criteria 2: Does the problem statement section clearly explain the importance of this specific work?
Evaluation Criteria 3: Does the problem statement section clearly summarize the previously proposed solutions and analyze their advantages and disadvantages?
Evaluation Criteria 4: Does the problem statement section effectively position this project within existing solutions, justifying its contribution or necessity?
Max Points: 2
Your Points: 2
Comments: Try to clarify the research gap better, describe the specific limitations in prior methods.

Methodology

Evaluation Criteria 1: Does the methodology section provides a detailed description of the used datasets?
Evaluation Criteria 2: Does the methodology section detail the experimental methods to be used?
Evaluation Criteria 3: Does it outline a complete plan of work sufficient for producing a finalized product?
Evaluation Criteria 4: Is the proposed methodology and work plan justified by references to previous work?
Max Points: 2
Your Points: 1
Comments: Add evaluation metrics and compare with current methods.

Work Plan

Note: May be integrated into other sections if explicitly addressed

Evaluation Criteria 1: Is the overall plan of work clearly defined?
Evaluation Criteria 2: Is the plan realistic given the time frame and resources?
Max Points: 2
Your Points: 2
Comments: Future work is clearly described.

Preliminary Results

Evaluation Criteria 1: Do the preliminary results demonstrate the viability of the proposed methods?
Evaluation Criteria 2: Do the preliminary results justify continuing with the chosen work plan? (i.e., do they suggest the plan is likely to succeed?)
Max Points: 2
Your Points: 2
Comments: The results are good, but requires additional refinement: try to improve judge consistency, expand validation and provide the comparison with existing approaches.

Reviewer Two

Report Structure

Evaluation Criteria: Does the report include all required sections: Abstract, Problem Statement, Methodology, and (Preliminary) Results?
Max Points: 1
Your Points: 1
Comments: The structure is clear, and all necessary sections are included.

Abstract

Evaluation Criteria 1: Does the abstract clearly state the significance of the work?
Evaluation Criteria 2: Does the abstract clearly reflect the area, in which the project results will be implemented?
Max Points: 1
Your Points: 1
Comments: Clear motivation; consider adding one headline number for impact.

Problem Statement

Evaluation Criteria 1: Does the problem statement section provide a comprehensive explanation of the project task?
Evaluation Criteria 2: Does the problem statement section clearly explain the importance of this specific work?
Evaluation Criteria 3: Does the problem statement section clearly summarize the previously proposed solutions and analyze their advantages and disadvantages?
Evaluation Criteria 4: Does the problem statement section effectively position this project within existing solutions, justifying its contribution or necessity?
Max Points: 2
Your Points: 1
Comments: The problem is explained well, but the gap in existing methods could be clearer.

Methodology

Evaluation Criteria 1: Does the methodology section provides a detailed description of the used datasets?
Evaluation Criteria 2: Does the methodology section detail the experimental methods to be used?
Evaluation Criteria 3: Does it outline a complete plan of work sufficient for producing a finalized product?
Evaluation Criteria 4: Is the proposed methodology and work plan justified by references to previous work?
Max Points: 2
Your Points: 1
Comments: The method is interesting, but more details about the data and prompts are needed to repeat the work.

Work Plan

Note: May be integrated into other sections if explicitly addressed

Evaluation Criteria 1: Is the overall plan of work clearly defined?
Evaluation Criteria 2: Is the plan realistic given the time frame and resources?
Max Points: 2
Your Points: 2
Comments: The work plan is clear and doable within the given time.

Preliminary Results

Evaluation Criteria 1: Do the preliminary results demonstrate the viability of the proposed methods?
Evaluation Criteria 2: Do the preliminary results justify continuing with the chosen work plan? (i.e., do they suggest the plan is likely to succeed?)
Max Points: 2
Your Points: 2
Comments: The results are promising; adding one more test would make them stronger.

Reviewer Three

Reviewer One

Report Structure

Evaluation Criteria: Does the report include all required sections: Abstract, Problem Statement, Methodology, and (Preliminary) Results?
Max Points: 1
Your Points: 1
Comments: The report successfully includes all necessary sections, adhering to the standard structure of an academic paper.

Abstract

Evaluation Criteria 1: Does the abstract clearly state the significance of the work?
Evaluation Criteria 2: Does the abstract clearly reflect the area, in which the project results will be implemented?
Max Points: 1
Your Points: 1
Comments: The abstract clearly outlines the significance of the work by stating that "Central banks face significant challenges in effectively communicating with the public, which is essential for influencing economic decision-making and fostering trust that can lead to lower inflation expectations". It further emphasizes that Large Language Models (LLM) offer a promising solution to enhance understanding and improve the effectiveness of central bank messages. The abstract clearly indicates the implementation area of the project results: "This project focuses on creating a synthetic focus group (SFG) of several LLM avatars of real people with different sentiments towards the central bank based on their interviews". This highlights the application in pre-testing communications for central banks.

Problem Statement

Evaluation Criteria 1: Does the problem statement section provide a comprehensive explanation of the project task?
Evaluation Criteria 2: Does the problem statement section clearly explain the importance of this specific work?
Evaluation Criteria 3: Does the problem statement section clearly summarize the previously proposed solutions and analyze their advantages and disadvantages?
Evaluation Criteria 4: Does the problem statement section effectively position this project within existing solutions, justifying its contribution or necessity?
Max Points: 2
Your Points: 2
Comments: The problem statement section comprehensively explains the project task, detailing the difficulties central banks face in public communication and the critical role of focus groups in understanding subjective perceptions. The objective of creating LLM-based synthetic focus groups is clearly articulated. The section thoroughly explains the work's importance, noting that higher trust in the Central Bank correlates with lower inflation expectations and improved economic decision-making. It highlights that traditional focus groups are logistically challenging and costly , which can be overcome by implementing LLM agents for communication pre-testing. The "Background Literature Review" section effectively summarizes previous digital focus group approaches, such as synchronous/asynchronous digital discussions, AI for transcription/sentiment analysis, and gamification. It critically analyzes their limitations, noting that they only partially reduce costs as they still rely on human respondents. The project is effectively positioned as a novel solution that replaces human respondents with LLM-based agents, aligning with the broader trend of AI-based digitalization of business processes. The necessity is justified by eliminating the drawbacks of traditional human focus groups and opening new avenues for regulatory policy pre-testing.

Methodology

Evaluation Criteria 1: Does the methodology section provides a detailed description of the used datasets?
Evaluation Criteria 2: Does the methodology section detail the experimental methods to be used?
Evaluation Criteria 3: Does it outline a complete plan of work sufficient for producing a finalized product?
Evaluation Criteria 4: Is the proposed methodology and work plan justified by references to previous work?
Max Points: 2
Your Points: 1
Comments: The methodology section mentions the use of "10 interviews is provided in txt format" and "the first page of the report" as test data. However, a more detailed description of the datasets, such as their structure, volume, or how they were collected (e.g., whether interviews represent diverse demographics or sentiments), is lacking. Providing anonymized examples of the interview content or the report excerpts would enhance clarity. The methodology section provides a highly detailed description of the experimental methods. It explains the two-step trust evaluation pipeline utilizing GigaChat models as agents with a one-shot prompting technique. The dual-phase architecture is clearly outlined: GigaChat-2-Max with a high temperature (0.9) for opinion generation, and the same model with a low temperature (0.1) as an "LLM-as-a-judge" for quantitative trust scores (0-1 scale). It also specifies that the full interview

text is added to the prompt for opinion formation. The "Objectives" section and "Extra work" section outline a comprehensive work plan, including LLM agent creation, calibration, synthetic focus group merging, discussion, and quality enhancement of the central bank's communication based on feedback. While the plan covers key steps, it lacks specific details on how "quality enhancement" will be achieved based on feedback, or what specific metrics beyond overall favorability will be used for evaluating the revised text. More granular detail on the iterative process and completion criteria would be beneficial. The proposed methodology and work plan are well-justified by references to previous work. For instance, the application of LLM agents as focus groups is supported by the advantages outlined in Zhang et al.. The ability of LLMs to simulate opinion dynamics and cognitive biases is referenced from Chuang et al.. The use of GigaChat models is justified by their convenient and scalable nature for prototyping.

Work Plan

Note: May be integrated into other sections if explicitly addressed

Evaluation Criteria 1: Is the overall plan of work clearly defined?
Evaluation Criteria 2: Is the plan realistic given the time frame and resources?
Max Points: 2
Your Points: 1.5
Comments: The overall plan of work is clearly defined in the "Objectives" section, encompassing literature review, LLM agent creation, calibration, SFG merging, discussions, and communication enhancement based on feedback. Additional tasks like innovative prompting strategies and sentiment-shift tracking are also noted. As a summer school project, the plan is ambitious but appears realistic for prototyping and obtaining preliminary results. However, the document does not specify concrete timelines beyond "SMILES Summer School Projects Proceedings", making a precise assessment of realism difficult. Resource availability (e.g., computational power for GigaChat-2-Max) is also not detailed. While acceptable for a prototype-focused summer school, more extensive research would require greater detail on timelines and resources.

Preliminary Results

Evaluation Criteria 1: Do the preliminary results demonstrate the viability of the proposed methods?
Evaluation Criteria 2: Do the preliminary results justify continuing with the chosen work plan? (i.e., do they suggest the plan is likely to succeed?)
Max Points: 2
Your Points: 2
Comments: The preliminary results demonstrate the viability of the proposed methods. The agents accurately reflect categorically positive opinions , and for wavering opinions, they were "quite similar to people". This indicates the model's capability to capture and replicate various sentiment types, supporting the method's feasibility. The preliminary results justify continuing with the chosen work plan. Despite some inconsistencies for "wavering" opinions and the observation that the model paid less attention to the test data than desired, the overall accuracy for categorical opinions and the general similarity to human responses suggest a promising approach. The authors acknowledge these limitations and outline future work on fine-tuning agents and improving response assessment mechanisms, indicating a clear path toward success.